

Solutions Exercise #7

Logistic Regression with R

Pasquale De Rosa

03/05/2022

Preliminaries (Problem 1)

Load the “Cars.txt” dataset from the ILIAS website.

```
cars <- read.table("Cars.txt", header = T)
summary(cars)
```

```
##      mpg      cylinders  displacement  horsepower      weight
## Min.   : 9.00   Min.   :3.000   Min.   : 68.0   Min.   : 46.0   Min.   :1613
## 1st Qu.:17.50   1st Qu.:4.000   1st Qu.:104.2   1st Qu.: 75.0   1st Qu.:2224
## Median :23.00   Median :4.000   Median :148.5   Median : 93.5   Median :2804
## Mean   :23.51   Mean   :5.455   Mean   :193.4   Mean   :104.5   Mean   :2970
## 3rd Qu.:29.00   3rd Qu.:8.000   3rd Qu.:262.0   3rd Qu.:126.0   3rd Qu.:3608
## Max.   :46.60   Max.   :8.000   Max.   :455.0   Max.   :230.0   Max.   :5140
##
##      acceleration      year      origin      name
## Min.   : 8.00   Min.   :70.00   Min.   :1.000   ford pinto      : 6
## 1st Qu.:13.82   1st Qu.:73.00   1st Qu.:1.000   amc matador     : 5
## Median :15.50   Median :76.00   Median :1.000   ford maverick   : 5
## Mean   :15.57   Mean   :76.01   Mean   :1.573   toyota corolla  : 5
## 3rd Qu.:17.18   3rd Qu.:79.00   3rd Qu.:2.000   amc gremlin     : 4
## Max.   :24.80   Max.   :82.00   Max.   :3.000   amc hornet      : 4
##                                     (Other)      :369
```

Problem 1: Consider the Cars dataset:

a. Build three different (generalized) linear regression models to predict mpg (at least one of them must be a multiple regression model).

Analyzing the dataset, we can assume that the model name is not useful if our aim is to predict the car system performance. Therefore, we can eliminate this variable.

```
cars_new <- cars[, c("mpg", "cylinders", "displacement", "horsepower", "weight",
                    "acceleration", "year", "origin")]
summary(cars_new)
```

```
##      mpg      cylinders  displacement  horsepower      weight
## Min.   : 9.00   Min.   :3.000   Min.   : 68.0   Min.   : 46.0   Min.   :1613
## 1st Qu.:17.50   1st Qu.:4.000   1st Qu.:104.2   1st Qu.: 75.0   1st Qu.:2224
## Median :23.00   Median :4.000   Median :148.5   Median : 93.5   Median :2804
## Mean   :23.51   Mean   :5.455   Mean   :193.4   Mean   :104.5   Mean   :2970
## 3rd Qu.:29.00   3rd Qu.:8.000   3rd Qu.:262.0   3rd Qu.:126.0   3rd Qu.:3608
## Max.   :46.60   Max.   :8.000   Max.   :455.0   Max.   :230.0   Max.   :5140
```

```
##
##      acceleration      year      origin      NA's      :6
##  Min.   : 8.00   Min.   :70.00   Min.   :1.000
## 1st Qu.:13.82   1st Qu.:73.00   1st Qu.:1.000
## Median :15.50   Median :76.00   Median :1.000
## Mean   :15.57   Mean   :76.01   Mean   :1.573
## 3rd Qu.:17.18   3rd Qu.:79.00   3rd Qu.:2.000
## Max.   :24.80   Max.   :82.00   Max.   :3.000
##
```

Before proceeding with our analyses, we can standardize our values in order to have all the predictors in the same measurement scale and speed up the model training:

```
to_remove <- which(is.na(cars_new$horsepower))
cars_new <- cars_new[-to_remove, ]
means_cars <- lapply(cars_new, mean)
sd_cars <- lapply(cars_new, sd)
cars_standardized <- (cars_new - means_cars) / sd_cars
attach(cars_standardized)
```

We can select the two most important predictors to build the first two generalized regression models. We first choose the one that minimizes the residual sum of squares (RSS):

```
predictors_cars <- names(cars_standardized)
predictors_cars <- predictors_cars[predictors_cars != "mpg"]
n_predictors_cars <- length(predictors_cars)
RSS_cars <- numeric(n_predictors_cars)

for (i in 1:n_predictors_cars) {
  cars_lm <- lm(mpg ~ cars_standardized[predictors_cars][, i],
               data=cars_standardized)
  RSS_cars[i] <- sum(cars_lm$residuals^2)
}

predictors_cars[which(RSS_cars== min(RSS_cars))]
```

```
## [1] "weight"
min(RSS_cars)
```

```
## [1] 120.1815
```

So weight is the best predictor, that minimizes the residual sum of squares. We can then find the second best one:

```
predictors_cars <- names(cars_standardized)
predictors_cars <- predictors_cars[predictors_cars != "mpg"]
predictors_cars <- predictors_cars[predictors_cars != "weight"]
n_predictors_cars <- length(predictors_cars)
RSS_cars <- numeric(n_predictors_cars)

for (i in 1:n_predictors_cars) {
  cars_lm <- lm(mpg ~ cars_standardized[predictors_cars][, i],
               data=cars_standardized)
  RSS_cars[i] <- sum(cars_lm$residuals^2)
}

predictors_cars[which(RSS_cars== min(RSS_cars))]
```

```
## [1] "displacement"
```

```
min(RSS_cars)
```

```
## [1] 137.5423
```

We found that weight and displacement are the two variables that minimize the RSS. We proceed with building our two single regression models:

```
weight_glm <- glm(mpg ~ weight, data=cars_standardized)
summary(weight_glm)
```

```
##
## Call:
## glm(formula = mpg ~ weight, data = cars_standardized)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.53409  -0.35305  -0.04303   0.27391   2.11651
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  2.663e-16  2.804e-02   0.00      1
## weight      -8.322e-01  2.807e-02  -29.64   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.3081577)
##
##      Null deviance: 391.00  on 391  degrees of freedom
## Residual deviance: 120.18  on 390  degrees of freedom
## AIC: 655
##
## Number of Fisher Scoring iterations: 2
```

```
displacement_glm <- glm(mpg ~ displacement, data=cars_standardized)
summary(displacement_glm)
```

```
##
## Call:
## glm(formula = mpg ~ displacement, data = cars_standardized)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.65497  -0.38748  -0.06433   0.30124   2.38473
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  5.584e-16  2.999e-02   0.00      1
## displacement -8.051e-01  3.003e-02  -26.81   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.3526726)
##
```

```
## Null deviance: 391.00 on 391 degrees of freedom
## Residual deviance: 137.54 on 390 degrees of freedom
## AIC: 707.89
##
## Number of Fisher Scoring iterations: 2
```

We can see that both model are statistically significant. Let's create a third glm, this time multiple:

```
cars_multiple_glm <- glm(mpg ~ ., data=cars_standardized)
summary(cars_multiple_glm)
```

```
##
## Call:
## glm(formula = mpg ~ ., data = cars_standardized)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.22873  -0.27630  -0.01498   0.23946   1.67334
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.046e-15  2.153e-02   0.000  1.00000
## cylinders    -1.078e-01  7.065e-02  -1.526  0.12780
## displacement  2.667e-01  1.008e-01   2.647  0.00844 **
## horsepower   -8.360e-02  6.799e-02  -1.230  0.21963
## weight       -7.046e-01  7.096e-02  -9.929 < 2e-16 ***
## acceleration  2.848e-02  3.494e-02   0.815  0.41548
## year         3.543e-01  2.406e-02  14.729 < 2e-16 ***
## origin       1.472e-01  2.871e-02   5.127 4.67e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.1817762)
##
## Null deviance: 391.000 on 391 degrees of freedom
## Residual deviance: 69.802 on 384 degrees of freedom
## AIC: 454.01
##
## Number of Fisher Scoring iterations: 2
```

Let's remove the non-significant variables (cylinders, horsepower and acceleration:

```
cars_multiple_glm <- glm(mpg ~ displacement + weight + year + origin,
                        data=cars_standardized)
summary(cars_multiple_glm)
```

```
##
## Call:
## glm(formula = mpg ~ displacement + weight + year + origin, data = cars_standardized)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.25691  -0.27071  -0.00497   0.22710   1.69232
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept)  1.128e-15  2.165e-02  0.000    1.000
## displacement 7.492e-02  6.393e-02  1.172    0.242
## weight      -7.156e-01  6.063e-02 -11.802 < 2e-16 ***
## year         3.641e-01  2.351e-02  15.486 < 2e-16 ***
## origin       1.265e-01  2.755e-02   4.593 5.92e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.1837824)
##
## Null deviance: 391.000 on 391 degrees of freedom
## Residual deviance: 71.124 on 387 degrees of freedom
## AIC: 455.37
##
## Number of Fisher Scoring iterations: 2
```

Now the displacement is not significant anymore. We will remove it to create our final model:

```
cars_multiple_glm <- glm(mpg ~ weight + year + origin, data=cars_standardized)
summary(cars_multiple_glm)
```

```
##
## Call:
## glm(formula = mpg ~ weight + year + origin, data = cars_standardized)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.27406  -0.26839  -0.00499   0.22108   1.70047
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.135e-15  2.166e-02   0.000      1
## weight      -6.523e-01  2.765e-02 -23.588 < 2e-16 ***
## year         3.573e-01  2.281e-02  15.668 < 2e-16 ***
## origin       1.187e-01  2.674e-02   4.439 1.18e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.1839593)
##
## Null deviance: 391.000 on 391 degrees of freedom
## Residual deviance: 71.376 on 388 degrees of freedom
## AIC: 454.76
##
## Number of Fisher Scoring iterations: 2
```

b. Perform 10-fold cross validation to estimate the test error of the models you built in a).

We can easily compare the three models performing a 10-fold cv using the `cv.glm()` function from the library `boot`:

```
library(boot)
set.seed(123)
weight_cv <- cv.glm(cars_standardized, weight_glm, K=10)
weight_cv$delta[1]
```

```
## [1] 0.3091372
```

```
displacement_cv <- cv.glm(cars_standardized, displacement_glm, K=10)
displacement_cv$delta[1]
```

```
## [1] 0.3539645
```

```
cars_multiple_cv <- cv.glm(cars_standardized, cars_multiple_glm, K=10)
cars_multiple_cv$delta[1]
```

```
## [1] 0.186302
```

We see that the multiple regression model is the best performing one with an average RMSE of ≈ 0.18 .

Preliminaries (Problem 2)

Load the “Cancer.txt” dataset from the ILIAS website.

```
cancer <- read.table("Cancer.txt", header = T)
summary(cancer)
```

```
##           ID           Diagnostic           Radius           Texture
## Min.      :    8670      B:357      Min.      : 6.981      Min.      : 9.71
## 1st Qu.:   869218      M:212      1st Qu.:11.700      1st Qu.:16.17
## Median :   906024                        Median :13.370      Median :18.84
## Mean      : 30371831                        Mean      :14.127      Mean      :19.29
## 3rd Qu.:   8813129                        3rd Qu.:15.780      3rd Qu.:21.80
## Max.      :911320502                        Max.      :28.110      Max.      :39.28
##      Perimeter           Area           Smooth           Compact
## Min.      : 43.79      Min.      : 143.5      Min.      :0.05263      Min.      :0.01938
## 1st Qu.: 75.17      1st Qu.: 420.3      1st Qu.:0.08637      1st Qu.:0.06492
## Median : 86.24      Median : 551.1      Median :0.09587      Median :0.09263
## Mean      : 91.97      Mean      : 654.9      Mean      :0.09636      Mean      :0.10434
## 3rd Qu.:104.10      3rd Qu.: 782.7      3rd Qu.:0.10530      3rd Qu.:0.13040
## Max.      :188.50      Max.      :2501.0      Max.      :0.16340      Max.      :0.34540
##      Concavity           Concave           Symmetry           Fractal
## Min.      :0.00000      Min.      :0.00000      Min.      :0.1060      Min.      :0.04996
## 1st Qu.:0.02956      1st Qu.:0.02031      1st Qu.:0.1619      1st Qu.:0.05770
## Median :0.06154      Median :0.03350      Median :0.1792      Median :0.06154
## Mean      :0.08880      Mean      :0.04892      Mean      :0.1812      Mean      :0.06280
## 3rd Qu.:0.13070      3rd Qu.:0.07400      3rd Qu.:0.1957      3rd Qu.:0.06612
## Max.      :0.42680      Max.      :0.20120      Max.      :0.3040      Max.      :0.09744
##      RadiusSE           TextureSE           PerimeterSE           AreaSE
## Min.      :0.1115      Min.      :0.3602      Min.      : 0.757      Min.      : 6.802
## 1st Qu.:0.2324      1st Qu.:0.8339      1st Qu.: 1.606      1st Qu.: 17.850
## Median :0.3242      Median :1.1080      Median : 2.287      Median : 24.530
## Mean      :0.4052      Mean      :1.2169      Mean      : 2.866      Mean      : 40.337
## 3rd Qu.:0.4789      3rd Qu.:1.4740      3rd Qu.: 3.357      3rd Qu.: 45.190
## Max.      :2.8730      Max.      :4.8850      Max.      :21.980      Max.      :542.200
##      SmoothSE           CompactSE           ConcavitySE           ConcaveSE
## Min.      :0.001713      Min.      :0.002252      Min.      :0.00000      Min.      :0.000000
## 1st Qu.:0.005169      1st Qu.:0.013080      1st Qu.:0.01509      1st Qu.:0.007638
## Median :0.006380      Median :0.020450      Median :0.02589      Median :0.010930
## Mean      :0.007041      Mean      :0.025478      Mean      :0.03189      Mean      :0.011796
## 3rd Qu.:0.008146      3rd Qu.:0.032450      3rd Qu.:0.04205      3rd Qu.:0.014710
## Max.      :0.031130      Max.      :0.135400      Max.      :0.39600      Max.      :0.052790
##      SymmetrySE           FractalSE           RadiusMax           TextureMax
```

```
## Min. :0.007882 Min. :0.0008948 Min. : 7.93 Min. :12.02
## 1st Qu.:0.015160 1st Qu.:0.0022480 1st Qu.:13.01 1st Qu.:21.08
## Median :0.018730 Median :0.0031870 Median :14.97 Median :25.41
## Mean :0.020542 Mean :0.0037949 Mean :16.27 Mean :25.68
## 3rd Qu.:0.023480 3rd Qu.:0.0045580 3rd Qu.:18.79 3rd Qu.:29.72
## Max. :0.078950 Max. :0.0298400 Max. :36.04 Max. :49.54
## PerimeterMax AreaMax SmoothMax CompactMax
## Min. : 50.41 Min. : 185.2 Min. :0.07117 Min. :0.02729
## 1st Qu.: 84.11 1st Qu.: 515.3 1st Qu.:0.11660 1st Qu.:0.14720
## Median : 97.66 Median : 686.5 Median :0.13130 Median :0.21190
## Mean :107.26 Mean : 880.6 Mean :0.13237 Mean :0.25427
## 3rd Qu.:125.40 3rd Qu.:1084.0 3rd Qu.:0.14600 3rd Qu.:0.33910
## Max. :251.20 Max. :4254.0 Max. :0.22260 Max. :1.05800
## ConcavityMax ConcaveMax SymmetryMax FractalMax
## Min. :0.0000 Min. :0.00000 Min. :0.1565 Min. :0.05504
## 1st Qu.:0.1145 1st Qu.:0.06493 1st Qu.:0.2504 1st Qu.:0.07146
## Median :0.2267 Median :0.09993 Median :0.2822 Median :0.08004
## Mean :0.2722 Mean :0.11461 Mean :0.2901 Mean :0.08395
## 3rd Qu.:0.3829 3rd Qu.:0.16140 3rd Qu.:0.3179 3rd Qu.:0.09208
## Max. :1.2520 Max. :0.29100 Max. :0.6638 Max. :0.20750
```

Problem 2: Apply the logistic regression to predict the category diagnosis and interpret the most important values of the model that you obtained. Can you estimate the error rate of your model?

We can get rid of the ID variable that cannot be used as a model predictor:

```
cancer_new <- cancer[,-1]
attach(cancer_new)
```

Let's create our logistic model to predict the variable Diagnostic using glm() with family=binomial as an argument:

```
cancer_logr <- glm(Diagnostic ~ ., data=cancer_new, family=binomial)
summary(cancer_logr)
```

```
##
## Call:
## glm(formula = Diagnostic ~ ., family = binomial, data = cancer_new)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -8.49   -8.49   -8.49    8.49    8.49
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -2.881e+06  2.816e+05 -10.233 < 2e-16 ***
## Radius       2.427e+06  2.693e+05   9.014 < 2e-16 ***
## Texture      1.958e+05  1.471e+04  13.313 < 2e-16 ***
## Perimeter    1.473e+06  2.464e+04  59.791 < 2e-16 ***
## Area        -1.301e+05  3.907e+03 -33.301 < 2e-16 ***
## Smooth      -1.525e+08  8.361e+06 -18.234 < 2e-16 ***
## Compact     -6.428e+06  3.213e+06  -2.001  0.04539 *
## Concavity    1.042e+06  1.408e+06   0.740  0.45959
## Concave     -1.716e+07  5.382e+06  -3.188  0.00143 **
```

```
## Symmetry      4.049e+07  7.772e+05  52.093 < 2e-16 ***
## Fractal      -4.233e+07  2.169e+06 -19.519 < 2e-16 ***
## RadiusSE     3.328e+07  1.169e+06  28.478 < 2e-16 ***
## TextureSE    6.368e+06  2.005e+05  31.763 < 2e-16 ***
## PerimeterSE  1.701e+06  4.720e+04  36.032 < 2e-16 ***
## AreaSE      -6.393e+05  1.835e+04 -34.840 < 2e-16 ***
## SmoothSE     7.492e+08  1.224e+07  61.213 < 2e-16 ***
## CompactSE   -1.773e+08  5.732e+06 -30.931 < 2e-16 ***
## ConcavitySE  1.529e+08  5.340e+06  28.624 < 2e-16 ***
## ConcaveSE   -1.260e+09  4.012e+07 -31.398 < 2e-16 ***
## SymmetrySE   2.890e+08  4.126e+06  70.055 < 2e-16 ***
## FractalSE    1.512e+09  6.597e+07  22.921 < 2e-16 ***
## RadiusMax   -6.130e+06  2.143e+05 -28.606 < 2e-16 ***
## TextureMax  -5.832e+05  2.437e+04 -23.935 < 2e-16 ***
## PerimeterMax -3.538e+05  1.219e+04 -29.023 < 2e-16 ***
## AreaMax     8.950e+04  2.741e+03  32.658 < 2e-16 ***
## SmoothMax   -2.161e+07  3.298e+06 -6.553 5.66e-11 ***
## CompactMax  8.986e+06  3.999e+05  22.470 < 2e-16 ***
## ConcavityMax -3.028e+07  1.523e+06 -19.875 < 2e-16 ***
## ConcaveMax  1.431e+08  5.471e+06  26.162 < 2e-16 ***
## SymmetryMax -2.474e+07  3.392e+05 -72.923 < 2e-16 ***
## FractalMax  -3.698e+07  5.340e+06 -6.926 4.33e-12 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 751.44 on 568 degrees of freedom
## Residual deviance: 32006.76 on 538 degrees of freedom
## AIC: 32069
##
## Number of Fisher Scoring iterations: 25
```

As we can see, the Concavity is not significant. We can then remove it and re-create our model:

```
cancer_new <- cancer_new[,-8]
cancer_logr <- glm(Diagnostic ~ ., data=cancer_new, family=binomial)
summary(cancer_logr)

##
## Call:
## glm(formula = Diagnostic ~ ., family = binomial, data = cancer_new)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
##    -8.49     0.00     0.00     0.00     8.49
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.217e+14  1.213e+08 -1003344 <2e-16 ***
## Radius      -2.927e+15  4.760e+07 -61502056 <2e-16 ***
## Texture      6.820e+13  2.257e+06  30219323 <2e-16 ***
## Perimeter    1.702e+14  6.892e+06  24690108 <2e-16 ***
## Area         1.624e+13  1.491e+05  108946317 <2e-16 ***
## Smooth       1.746e+16  5.657e+08  30867575 <2e-16 ***
```



```
## Compact      -1.771e+16  3.776e+08 -46912470 <2e-16 ***
## Concave       2.429e+16  4.045e+08  60040200 <2e-16 ***
## Symmetry     -7.072e+15  2.109e+08 -33536226 <2e-16 ***
## Fractal      -7.459e+15  1.580e+09 -4721581  <2e-16 ***
## RadiusSE      6.869e+14  8.812e+07  7794767  <2e-16 ***
## TextureSE     2.922e+12  1.046e+07  279437   <2e-16 ***
## PerimeterSE  -4.643e+14  1.163e+07 -39917322 <2e-16 ***
## AreaSE        3.466e+13  3.945e+05  87868453 <2e-16 ***
## SmoothSE     -1.701e+16  1.877e+09 -9062133  <2e-16 ***
## CompactSE     3.244e+16  6.099e+08  53194503 <2e-16 ***
## ConcavitySE  -2.676e+16  3.234e+08 -82738584 <2e-16 ***
## ConcaveSE     1.175e+17  1.534e+09  76621241 <2e-16 ***
## SymmetrySE    -2.541e+16  7.748e+08 -32789711 <2e-16 ***
## FractalSE    -2.677e+17  3.314e+09 -80763694 <2e-16 ***
## RadiusMax     7.154e+14  1.647e+07  43440178 <2e-16 ***
## TextureMax    -6.198e+12  1.974e+06 -3139500  <2e-16 ***
## PerimeterMax  3.523e+13  1.686e+06  20891171 <2e-16 ***
## AreaMax      -5.185e+12  9.069e+04 -57177752 <2e-16 ***
## SmoothMax     7.875e+14  4.076e+08  1932247  <2e-16 ***
## CompactMax   -2.923e+15  1.086e+08 -26910623 <2e-16 ***
## ConcavityMax  3.729e+15  6.372e+07  58528461 <2e-16 ***
## ConcaveMax   -5.255e+15  2.541e+08 -20685733 <2e-16 ***
## SymmetryMax   6.714e+15  1.404e+08  47822073 <2e-16 ***
## FractalMax    2.888e+16  6.748e+08  42795823 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 751.44  on 568  degrees of freedom
## Residual deviance: 865.05  on 539  degrees of freedom
## AIC: 925.05
##
## Number of Fisher Scoring iterations: 25
```

We can observe that the computation was done after 25 iterations. The coefficients are now all statistically significant, with values either positive or negative (if $\beta_n > 0$, for higher values of x_n increases the probability that $y = 1$, and vice-versa). To estimate the error rate of our model, let's recall the formula of the (multiple) logistic regression:

$$\Pi = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n)}}$$

Where Π is the probability of y to belong to a binary class and assumes values between $[0, 1]$. Since the output is binary, the error for each value of y will be:

$$\epsilon = \begin{cases} -\pi, & \text{if } y = 0 \\ 1 - \pi, & \text{if } y = 1 \end{cases}$$

Using the `cv.glm()` function, we can perform the 10-fold cross-validation and estimate the error rate (averaged on the 10 folds) of our logistic model:

```
cancer_cv <- cv.glm(cancer_new, cancer_logr, K=10)
cancer_cv$delta[1]
```

```
## [1] 0.05448154
```

So our model has an average error rate on the 10 folds of $\approx 5\%$, that is fairly low. We can then consider this

logistic model as a good predictor for the cancer diagnosis.